

Senior Software Engineer Coding Challenge: CEX-DEX Arbitrage Bot

Overview

Build a real-time arbitrage detection system in Go that monitors price discrepancies between a centralized exchange (CEX) orderbook and Uniswap V3 DEX pricing for the ETH-USDC trading pair.

Challenge Requirements

Core Functionality

1. CEX Orderbook Integration

- Integrate with Binance API to fetch ETH-USDC orderbook snapshots
- Snapshot the orderbook on every Ethereum block (~12 seconds)
- Calculate the effective execution price for configurable trade sizes (e.g., 1 ETH, 10 ETH, 100 ETH)

2. Ethereum Block Streaming

- **WebSocket Connection:** Establish a WebSocket connection to an Ethereum node (Infura, Alchemy, or local node)
- Subscribe to new block headers using `eth_subscribe` with `newHeads`
- **Connection Management:** Implement robust reconnection logic
 - Detect connection drops (ping/pong, timeout detection)
 - Automatic reconnection with exponential backoff
 - Resume from last known block to avoid gaps
 - Handle graceful degradation if WebSocket unavailable (fallback to polling)
- Process each block atomically to trigger orderbook snapshots

3. DEX Price Integration

- Connect to an Ethereum node (can use Infura/Alchemy)
- Query Uniswap V3 ETH-USDC pool for equivalent pricing
- Use the QuoterV2 contract or simulate swap calculations
- Fetch data synchronized with each new Ethereum block

Understanding Uniswap V3 Pools

Uniswap V3 uses a concentrated liquidity model where:

- Liquidity providers allocate capital within specific price ranges (ticks)
- The pool's current price is determined by the ratio of reserves and the current tick
- A swap moves through sequential ticks, consuming liquidity at each price level
- Pool state (`liquidity`, `sqrtPriceX96`, `tick`) changes with every swap transaction

Why Requote on Every Block:

- While you *could* use `eth_subscribe` with `logs` to watch for `Swap` events on the pool contract, requerying on every block is preferred because:
 1. **Atomic Consistency:** You get a consistent snapshot of both CEX and DEX state at a specific block height
 2. **No Event Gaps:** Log subscriptions can miss events during reconnections or node issues
 3. **State Changes Beyond Swaps:** Pool state can change through liquidity additions/removals, not just swaps
 4. **Arbitrage Window:** Knowing the price at each block boundary helps you understand the exact window of opportunity
 5. **Simpler Architecture:** Block-driven polling is more predictable than event-driven for this use case

For this challenge, implement the block-driven approach with requerying, but be prepared to discuss trade-offs with event-driven approaches.

3. Arbitrage Detection

- Compare CEX vs DEX pricing accounting for:
 - Slippage on both sides
 - Gas costs (estimate for DEX operations)
 - Trading fees
- Detect profitable arbitrage opportunities
- Print detailed arbitrage information including:
 - Direction (CEX→DEX or DEX→CEX)
 - Expected profit in USD and percentage
 - Execution steps required
 - Risk factors

Senior Engineer Requirements

Implement the following to demonstrate senior-level engineering skills:

1. Caching Strategy

- **Requirement:** Implement a multi-layered caching system
 - L1: In-memory cache for frequently accessed data (recent blocks, pool states)
 - L2: Consider cache invalidation strategies
 - Cache Uniswap pool state, gas price estimates, and historical orderbook data
- **Considerations:**
 - Thread-safe cache access
 - TTL-based expiration

- Cache warming strategies
- Memory bounds

2. WebSocket Management & Connection Resiliency

- **Requirement:** Implement production-grade WebSocket handling
 - Maintain persistent WebSocket connection to Ethereum node
 - Implement heartbeat/ping mechanism to detect stale connections
 - Automatic reconnection with exponential backoff and jitter
 - Track last processed block to resume without gaps after reconnection
 - Handle edge cases: late blocks, missed blocks, connection during sync
 - Proper cleanup of connections and subscriptions

3. Concurrency & Performance

- **Requirement:** Design a concurrent architecture using Go patterns
 - Use goroutines and channels effectively
 - Implement worker pools for API calls
 - Handle backpressure when data arrives faster than processing
 - Avoid race conditions
- **Bonus:** Implement graceful shutdown with context cancellation

4. Rate Limiting & Resiliency

- **Requirement:** Implement production-ready error handling
 - Rate limiting for external API calls (CEX API, RPC endpoints)
 - Exponential backoff with jitter for retries
 - Circuit breaker pattern for failing services
 - Metrics/observability hooks (even if just structured logging)

5. Configuration & Extensibility

- **Requirement:** Make the system configurable and extensible
 - Support multiple trading pairs (design for ETH-USDC, but make it extensible)
 - Configurable trade sizes to check
 - Pluggable exchange adapters (interface-based design)
 - Configuration via file (YAML/JSON) or environment variables

6. Data Modeling & Architecture

- **Requirement:** Demonstrate clean architecture principles
 - Clear separation of concerns (data layer, business logic, external integrations)
 - Well-defined interfaces between components
 - Proper error types and handling
 - Consider how you'd test this system (testability)

Expected Output

When an arbitrage opportunity is detected, print output similar to:

```
==== ARBITRAGE OPPORTUNITY DETECTED ====
Block Number: 18234567
Timestamp: 2024-01-15 14:23:45 UTC
Direction: CEX → DEX (Buy on Binance, Sell on Uniswap)

Trade Size: 10.0 ETH
CEX Price: $2,245.30 (effective with slippage)
DEX Price: $2,267.80 (effective with slippage)
Price Difference: $22.50 per ETH (1.00%)

Estimated Profit: $225.00 (before gas and fees)

Execution Steps:
1. Buy 10.0 ETH on Binance at average price $2,245.30
   - Place market order or limit orders consuming top bid/ask levels
   - Required capital: ~$22,453.00 USDC
2. Transfer ETH to trading wallet
3. Execute Uniswap V3 swap: 10.0 ETH → USDC
   - Pool: 0x88e6A0c2dDD26FEEb64F039a2c41296FcB3f5640
   - Expected output: ~22,678 USDC (after 0.3% pool fee)
```

Deliverables

1. Source Code

- Well-structured Go modules
- Clear package organization
- Idiomatic Go code

2. Documentation

- README with setup instructions
- Architecture diagram or description
- API documentation for key interfaces
- Configuration examples

3. Testing (Bonus)

- Unit tests for critical business logic
- Mock implementations for external services
- Integration test examples

4. Discussion Points (Be prepared to discuss)

- How would you deploy this in production?
- How would you monitor this system?
- What are the scaling bottlenecks?
- How would you handle chain reorgs?
- Security considerations for handling private keys (if extending to actual trading)
- How would you extend this to support multiple DEXes or CEXes?

Technical Constraints

- **Language:** Go 1.21+
- **External Dependencies:** Allowed (but be judicious)
 - Ethereum clients: go-ethereum (geth)
 - HTTP clients: standard library or popular choices
 - Configuration: viper, yaml, etc.
- **RPC Access:** Use public endpoints (Infura, Alchemy, or similar)
- **No Trading Required:** Detection only, no actual transaction submission

Evaluation Criteria

We evaluate candidates across four key dimensions:

1. Code Quality

- Idiomatic Go patterns and conventions
- Clean architecture and code organization
- Comprehensive error handling
- Code readability and maintainability

2. System Design

- Concurrency design and patterns
- WebSocket connection management and reconnection logic
- Caching strategy and implementation
- Extensibility and separation of concerns

3. Production Readiness

- Error handling & resiliency patterns
- Connection failure recovery
- Configuration management
- Logging/observability
- Resource management and cleanup

4. DeFi & Domain Understanding

- Accurate pricing calculations
- Understanding of Uniswap V3 mechanics
- Gas cost modeling
- Understanding of arbitrage mechanics and execution
- Consideration of real-world constraints and edge cases

Time Expectation

This challenge is designed to take **4-6 hours** for a senior engineer. Focus on:

- Core functionality working correctly
- Demonstrating 2-3 senior engineering patterns well (rather than all of them superficially)
- Clean, readable code with good architecture

We value quality over completeness. It's better to implement fewer features with production-quality code than to rush through all requirements. Document any shortcuts or TODOs for what you'd improve with more time.

Hints & Resources

WebSocket Ethereum Connection

- **Libraries:**

- gorilla/websocket - Most popular WebSocket client: go get github.com/gorilla/websocket
- nhooyr.io/websocket - Modern alternative with better context support

- **Connection Setup:**

```
// Example WebSocket URL for Infura
wsURL := "wss://mainnet.infura.io/ws/v3/YOUR_API_KEY"

// Subscribe to new block headers
subscribeMsg := `{"jsonrpc":"2.0","id":1,"method":"eth_subscribe","params":["newHeads"]}`
```

- **Message Handling:**

- Subscription confirmation: {"jsonrpc":"2.0","id":1,"result":"0x..."}
- Block notifications: {"jsonrpc":"2.0","method":"eth_subscription","params":{"subscription":"0x...","result":{...}}}
- Implement timeout detection (if no message in 30s, reconnect)

Getting Uniswap V3 Pool Quotes

There are two main approaches to get price quotes from Uniswap V3:

Option 1: Using QuoterV2 Contract (Recommended for Beginners)

The QuoterV2 contract simulates swaps and returns expected output amounts without executing trades.

Contract Details:

- QuoterV2 Address: 0xb27308f9F90D607463bb33eA1BeBb41C27CE5AB6
- Method: quoteExactInputSingle
- [Etherscan](#)

Go Implementation:

```
// 1. Install go-ethereum
// go get github.com/ethereum/go-ethereum

// 2. Generate Go bindings from ABI (or use pre-generated)
// abigen --abi quoterv2.abi --pkg uniswap --type QuoterV2 --out quoterv2.go

// 3. Call the contract
import (
    "github.com/ethereum/go-ethereum/ethclient"
    "github.com/ethereum/go-ethereum/common"
    "math/big"
)

client, _ := ethclient.Dial("https://mainnet.infura.io/v3/YOUR_KEY")
quoterAddr := common.HexToAddress("0xb27308f9F90D607463bb33eA1BeBb41C27CE5AB6")

// QuoteExactInputSingle parameters:
// - tokenIn: WETH address (0xC02aaA39b223FE8D0A0e5C4F27eAD9083C756Cc2)
// - tokenOut: USDC address (0xA0b86991c6218b36c1d19D4a2e9Eb0cE3606eB48)
// - fee: 3000 (0.3%)
// - amountIn: amount in wei (e.g., 1 ETH = 1e18)
// - sqrtPriceLimitX96: 0 (no limit)

params := QuoteExactInputSingleParams{
    TokenIn:        common.HexToAddress("0xC02aaA39b223FE8D0A0e5C4F27eAD9083C756Cc2"),
    TokenOut:       common.HexToAddress("0xA0b86991c6218b36c1d19D4a2e9Eb0cE3606eB48"),
    Fee:            big.NewInt(3000),
    AmountIn:       big.NewInt(1e18), // 1 ETH
    SqrtPriceLimitX96: big.NewInt(0),
}

// Use eth_call to simulate (no gas cost)
amountOut, sqrtPriceX96After, _, gasEstimate, _ := quoter.QuoteExactInputSingle(params)
// amountOut is in USDC units (6 decimals), e.g., 2250000000 = 2250 USDC
```

Without Code Generation (using ABI directly):

```
import (
    "github.com/ethereum/go-ethereum/accounts/abi"
    "github.com/ethereum/go-ethereum/common"
    "strings"
)

// Load ABI
quoterABI := `[{"inputs": [{"components": [{"internalType": "address", "name": "tokenIn", "type": "address"}, ...]}], "name": "quoteExactInputSingle"}]`
```

parsedABI, _ := abi.JSON(strings.NewReader(quoterABI))

// Pack the call data

```
data, _ := parsedABI.Pack("quoteExactInputSingle",
    common.HexToAddress("0xC02aa39b223FE8D0A0e5C4F27eAD9083C756Cc2"), // WETH
    common.HexToAddress("0xA0b86991c6218b36c1d19D4a2e9Eb0cE3606eB48"), // USDC
    big.NewInt(3000), // fee
    big.NewInt(1e18), // amountIn
    big.NewInt(0), // sqrtPriceLimitX96
)
```

// Make eth_call

```
msg := ethereum.CallMsg{
    To:   &quotercaller,
    Data: data,
}
result, _ := client.CallContract(context.Background(), msg, nil)
```

// Unpack result

```
var out []interface{}
parsedABI.UnpackIntoInterface(&out, "quoteExactInputSingle", result)
amountOut := out[0].(*big.Int)
```

Option 2: Direct Pool State Calculation (Advanced)

Read pool state directly and calculate swap amounts using Uniswap's math.

Pool Contract:

- ETH-USDC 0.3% Pool: 0x88e6A0c2dDD26FEEb64F039a2c41296FcB3f5640
- Method: slot0() returns current price and tick
- Also need: liquidity(), tickBitmap(), ticks()

Key Functions:

```
// Call slot0() to get current state
// Returns: sqrtPriceX96, tick, observationIndex, observationCardinality, ...

// Then implement Uniswap V3 math:
// - Convert sqrtPriceX96 to actual price
// - Walk through ticks to calculate swap output
// - Account for liquidity at each tick
```

Libraries:

- github.com/ethereum/go-ethereum - Core Ethereum client
- github.com/daoleno/uniswap-sdk-core - Uniswap SDK port (partial)
- github.com/daoleno/uniswapv3-sdk - V3 SDK port (useful for math)

Note: Option 2 requires implementing or porting Uniswap's tick math. For this challenge, **Option 1 (QuoterV2) is recommended** unless you want to demonstrate deep DeFi knowledge.

Important Addresses & Token Decimals

- **WETH:** 0xC02aa39b223FE8D0A0e5C4F27eAD9083C756Cc2 (18 decimals)
- **USDC:** 0xA0b86991c6218b36c1d19D4a2e9Eb0cE3606eB48 (6 decimals)
- **ETH-USDC Pool (0.3%):** 0x88e6A0c2dDD26FEEb64F039a2c41296FcB3f5640

- QuoterV2: 0xb27308f9F90D607463bb33eA1BeBb41C27CE5AB6

Understanding Token Decimals

ERC-20 tokens store amounts as integers without decimal points. The `decimals` property tells you how to interpret these integers.

How Decimals Work:

- **ETH/WETH (18 decimals)**: The smallest unit is 1 wei = 10^{-18} ETH
 - 1 ETH = 1,000,000,000,000,000,000 wei = 1×10^{18} wei
 - 0.5 ETH = 500,000,000,000,000,000 wei
 - When you see `big.NewInt(1e18)` in Go, that's 1 ETH
 - **USDC (6 decimals)**: The smallest unit is 1 micro-USDC = 10^{-6} USDC
 - 1 USDC = 1,000,000 micro-USDC = 1×10^6
 - 2,250 USDC = 2,250,000,000
 - 0.50 USDC = 500,000

Example Conversions in Go:

```
import (
    "math/big"
    "fmt"
)

// ETH (18 decimals)
oneETH := new(big.Int).Exp(big.NewInt(10), big.NewInt(18), nil)
fmt.Println(oneETH) // 1000000000000000000

// 10.5 ETH = 10.5 × 10^18
tenPointFiveETH := new(big.Int).Mul(big.NewInt(105), new(big.Int).Exp(big.NewInt(10), big.NewInt(17), nil))
// or: 1050000000000000000

// USDC (6 decimals)
oneUSDC := new(big.Int).Exp(big.NewInt(10), big.NewInt(6), nil)
fmt.Println(oneUSDC) // 100000

// 2,250.50 USDC = 2250.50 × 10^6 = 2,250,500,000
usdcAmount := big.NewInt(2250500000)
```

Converting from Smart Contract Response

```
// QuoterV2 returns amountOut in raw token units
amountOut := big.NewInt(2250500000) // USDC has 6 decimals

// Convert to human-readable format
divisor := new(big.Int).Exp(big.NewInt(10), big.NewInt(6), nil)
usdcFloat := new(big.Float).Quo(
    new(big.Float).SetInt(amountOut),
    new(big.Float).SetInt(divisor),
)
fmt.Printf("%.2f USDC\n", usdcFloat) // 2250.50 USDC
```

Why This Matters for Arbitrage:

- When calling QuoterV2 with 1 ETH input, pass `10000000000000000000` (1×10^{18})
 - The output will be in USDC units (6 decimals), e.g., `225000000` = 2,250 USDC
 - Binance prices are already in decimal format (strings like "2245.30")
 - You need to normalize both to compare: divide smart contract values by 10^6 decimals

Common Mistakes to Avoid:

- Comparing 2250000000 (raw USDC) directly with 2245.30 (Binance price)
 - Convert both to same format: $2250000000 / 10^6 = 2250.00$ vs 2245.30
 - Using float64 for large token amounts (precision loss)
 - Use big.Int for calculations, convert to big.Float only for display

Ethereum RPC

- **Block subscription:** `eth_subscribe` with `newHeads` parameter
 - **Contract calls:** `eth_call` (doesn't cost gas, read-only)

- **Gas estimates:** eth_gasPrice or eth_feeHistory
- **Rate limits:** Public endpoints (Infura/Alchemy) typically allow 100k requests/day on free tier

Binance API

- **Binance Spot API Documentation:** <https://binance-docs.github.io/apidocs/spot/en/#order-book>
 - **Orderbook Endpoint:** GET <https://api.binance.com/api/v3/depth>
 - Parameter: symbol=ETHUSDC
 - Parameter: limit=100 (or 500, 1000 for deeper orderbook)
 - **Example:** <https://api.binance.com/api/v3/depth?symbol=ETHUSDC&limit=100>
 - **Response Format:**

```
{
  "lastUpdateId": 1234567890,
  "bids": [["2245.30", "1.5"], ["2245.20", "2.3"], ...],
  "asks": [["2245.50", "1.2"], ["2245.60", "0.8"], ...]
}
```

- **Rate Limit:** 1200 requests per minute (no API key needed for public endpoints)
- **Note:** Prices are strings, amounts are strings (parse carefully)

Go Libraries Summary

```
# Essential
go get github.com/ethereum/go-ethereum
go get github.com/gorilla/websocket

# Optional but helpful
go get github.com/daoleno/uniswapv3-sdk # For Uniswap math utilities
go get github.com/shopspring/decimal      # For precise decimal arithmetic
```

ABI Files

You'll need the ABI for QuoterV2. You can:

1. Download from Etherscan: <https://etherscan.io/address/0xb27308f9F90D607463bb33eA1BeBb41C27CE5AB6#code>
2. Use abigen tool (comes with go-ethereum) to generate Go bindings
3. Or manually parse and call using accounts/abi package (shown above)

Questions?

Think about edge cases and trade-offs. We value engineers who:

- Ask clarifying questions
- Consider production implications
- Write maintainable code
- Think about observability and debugging

Good luck! We're excited to see your approach to this problem.